

What is File System?

- It is responsible for storing information on disk and retrieving and updating this information.
- Example :
 - FAT16, FAT32, NTFS
 - ext2, ext3
- In Linux everything is **file**.

Type of File System

- Network File System
 - NFS
 - SMB
- Disk File System
 - ext2
 - ext3
 - FAT32
 - NTFS

Network File System

- Network File System are physically somewhere else, but appear as if they are mounted on one computer.
- NFS
 - It was developed by Sun.
- SMB
 - It was developed by Microsoft.

Disk File System

- Disk File System are what you will find on a physical device, such as hard drive in a computer.

ext2 File System

- It has been the standard File System for Linux.
- The original **Ext**ended File System was named **ext**.
- The ext2 File System can accommodate:
 - Files as large as 2GB
 - Directories as large as 2TB
 - Max. file name length of 255 characters.

Hard Disk Partitions

- Disk divided into partitions
 - Definition of partition: group of adjacent cylinders
 - 1 file system may reside in a single partition
- BIOS defines boot sector to be head 0, cylinder 0, sector 1
- Master Boot Record (MBR) – used to boot computer
- Partition table
 - Start and end addresses of each partition
 - One partition is marked as active
- When computer boots up:
 - BIOS executes MBR
 - MBR locates active partition and executes boot block

Hard Disk Partitions

- Each partition:
 - Starts with **boot block** even if it does not contain bootable OS
 - **Superblock** – key parameters about file system in partition such as magic # identifying file system type, number of blocks
 - Bitmap or linked-list of free blocks – free space management
 - **Inodes** – array of these (one per file)
 - Root directory of partition
 - Files and directories at the end

ext2 Structure

- A file in the ext2 File System begins with the inode.
- inode
 - Each file has an inode structure that is identified by an i-number.
 - The inode contains the information required to access the file.
 - It doesn't contain file name.
 - Inodes store information on files, such as
 - user and group ownership,
 - access mode (read, write, execute permissions)
 - and type of file.


Physical Structure on the Disk

- Boot Block: information needs to boot the system
- Super Block: File System Specifications
 - Size
 - Max. number of files
 - Free blocks
 - Free inodes
- inode List
- Block List: The files data

ext3 File System

- It is as same as ext2.
- It is a **journaling** File System for Linux.
- In a journaling system, metadata is written to a journal on the disk before it is actually used to modify the file.

File System Structure



The Root Directory

- /bin
- /boot
- /dev
- /etc
- /home
- /initrd
- /lib
- /lost+found
- /media
- /mnt
- /opt
- /proc
- /root
- /sbin
- /usr
- /srv
- /tmp
- /var

/bin

- Hold the most commonly used essential user programs
 - login
 - Shells (bash, ksh, csh)
 - File manipulation utilities (cp, mv, rm, ln, tar)
 - Editors (ed, vi)
 - File system utilities (dd, df, mount, umount, sync)
 - System utilities (uname, hostname, arch)
 - GNU utilities like gzip and gunzip

/bin detail

- cat → Utility to concatenate files to standard output
- kill → Utility to send signals to processes
- chmod → Utility to change file access permissions
- chown → Utility to change file owner and group
- cp → Utility to copy files and directories
- ls → Utility to list directory contents
- mkdir → Utility to make directories
- mv → Utility to move/rename files
- pwd → Utility to print name of current working directory
- echo → Utility to display a line of text
- rm → Utility to remove files or directories

/sbin

- Hold essential maintenance or system programs such as the following:
 - fsck
 - Fdisk
 - Mkfs
 - Shutdown
 - Lilo
 - Init
- The main difference between the programs stored in /bin and /sbin is that the programs in /sbin are executable only by **root**.

/etc

- Store the **systemwide configuration** files required by many programs.
 - passwd
 - shadow
 - fstab
 - hosts
 - lilo.conf
 - ...

/home and /root

- The /home directory is where all the home directories for all the users on a system are stored.
- The /root directory is where all the home directories for root user on a system are stored.

/dev

- The special files representing hardware are kept in it.
 - /dev/hda1
 - /dev/ttyS0
 - /dev/mouse
 - /dev/fd0
 - ...

/tmp and /var

- The /tmp and /var directories are used to hold temporary files or files with constantly varying content.
- The /tmp directory is usually a dumping ground for files that only need to be used briefly and can afford to be deleted at any time.
- The /var directory is a bit more structured than /tmp and usually looks something like the following:
 - /var/log
 - /var/spool
 - /var/named
 - ...

/usr

- Most programs and files directly relating to users of the system are stored.
- It is in some ways a mini version of the / directory.
 - /usr/bin
 - /usr/sbin
 - /usr/spool
 - ...

/proc

- It is a virtual File System
- A special File System provided by the kernel as a way of providing information about the system to user programs.
- The main tasks of proc File System is to provide information about the kernel and processes.
- runtime system information (e.g. system memory, devices mounted, hardware configuration, etc).

Other directories

- /mnt
 - removable media such as CD-ROM, floppy and ... are mounted.
 - /mnt/floppy
 - /mnt/cdrom
- /boot
 - Image to boot system
- /lost+found
 - Used by fsck

/lib

- Contains kernel modules and those shared library images (the C programming code library) needed to boot the system and run the commands in the root filesystem,
 - ie. by binaries in /bin and /sbin
- Windows equivalent to a shared library would be a DLL (dynamically linked library) file

Mounting File System

- The Linux File System makes it appear as if all the File System are local and mounted somewhere on the root File System.
- File System are mounted with the **mount** command.
 - `mount -t type source mount_point`
- To unmount a File System, the **umount** command is used.
 - `umount /dev/<device name> or mount_point`

Mounting Automatically with fstab

- This file lists all the partitions that need to be mounted at boot time and the directory where they need to be mounted.
- Along with that information, you can pass parameters to the mount command.
- /etc/fstab
 - Which devices to be mounted
 - What kinds of File Systems they contain
 - At what point in the File System the mount takes place
 - ...

Partition Table

- MBR (Master Boot Record)
 - The first sector
 - 512 bytes (**446 bytes**:boot loader such as LILO or GRUB, **64 bytes**:partition table, **2 bytes**:special code).
- The partition table has enough room for **four** partitions.
 - One of the four can be used as an extended partition.

Partitions

- Primary-Master
 - /dev/hda
- Primary-Slave
 - /dev/hdb
- Secondary-Master
 - /dev/hdc
- Secondary-Slave
 - /dev/hdd
- Swap Partition
 - Used to implement virtual memory

Creating File System

- Once a disk has been partitioned for a specific File System, it is necessary to create a File System on it.
- The first process in the DOS world is known as **formatting**.
- In the UNIX world is known as creating a File System.

Create File System Commands

- mkfs or mke2fs
 - Make a new ext2 File System.
- mk3fs
 - Make a new ext3 File System.
- mkdosfs
 - Make DOS File System without owning any Microsoft software.

Formatting Linux Filesystem

- **Step #1 Create the new filesystem with following command (first login in as a root user)**
 - `mkfs.ext3 /dev/sda5`
- **Step # 2: Create mount point directory for the file system**
 - `# mkdir /datadisk1`
- **Step # 3: Mount the new file system**
 - `# mount /dev/sda5 /datadisk1`
- **Step # 4: Finally make sure file system /dev/hda5 automatically mounted at /datadisk1 mount point after system reboots.**
 - `vi /etc/fstab`
 - Add/append following entry to file:
 - `/dev/sda5 /datadisk1 ext3 defaults 0 2`

The shell

- The **\$ prompt** that you see when you first log in is displayed by a **shell** (usually **bash** is Linux default shell).
- The Shell executes programs.
 - User types command
 - Shell reads command (read from input) and translates it to the operating system.
 - **Shell types: Bash, csh, ksh, sh**
- The `[username@server current_directory]$` signifies that
 - this console is being used by user `username`,
 - and the host-name is `server`.
 - The second word is the current working directory

Commands

- General command syntax format:

`$ command -options arguments`

- Example:

➤ `$ clear`

➤ `$ cd /usr/src/linux`

➤ `$ wc -w file1` (number of words in file1)

➤ `$ wc -c file1` (number of characters in file1)

➤ `$ wc -l file1` (number of lines in file1)

➤ `$ cat file1 file2 file3`

➤ `$ ls -al`

Wild-cards

- `$ ls -l *.c`
- `$ ls [abc]*`
- `$ ls ?a*`

- `*` means 'match any number of characters'.
 - For example, `chap*` matches: *chap01*, *chapa*, *chap_end*, and also *chap*.
 - If you just give `*` (nothing else), it matches every file.
- `?` means 'match any single character'.
 - For example, `chap?` matches: *chapa* and *chap1*, but not *chap01* and *chap*.

Wild-cards

- [...] means 'match any one characters between the brackets'. A range of characters may be specified by separating a pair of characters by a dash.
 - For example, `chap[abc]` matches: *chapa* and *chapc*, but not *chap1* and *chapab*.
 - `[A-Za-z]*` matches with any word whose first element is a character
 - `[!abc]` matches with any one characters different from a,b and c

Control characters

- Interrupting
 - **^C** interrupts.
 - ✓ Exits the program and returns you to the command-line prompt.
 - **^Z** suspends.
 - ✓ Stops the program and puts it in the background. Type *fg* to restart it.
 - **^D** end of file.
- Stop/Start Scrolling: if what you want to see scrolls off the bottom of the screen, you may prevent this by sending a "stop" signal (**^S**) to the host. Send a "start" signal to resume (**^Q**).
 - **^S** stops scrolling
 - **^Q** resume scrolling

Linux Help Manual

- The `man` command displays help manual for selected command:

`man command_name`

- `$ man man` or `$ man -help` help for the `man` command itself!!

Example: the command `cat`

```
[rinaldi@homelinux rinaldi]$ man cat
```

CAT(1) **User Commands** **CAT(1)**

NAME

cat - concatenate files and print on the standard output

SYNOPSIS

cat [OPTION] [FILE]...

DESCRIPTION

Concatenate FILE(s), or standard input, to standard output.

-A, --show-all
equivalent to -vET

-b, --number-nonblank
number nonblank output lines

-e equivalent to -vE

-E, --show-ends
display \$ at end of each line

-n, --number
number all output lines

- s, --squeeze-blank**
never more than one single blank line
- t** equivalent to **-vT**
- T, --show-tabs**
display TAB characters as **^I**
- u** (ignored)
- v, --show-nonprinting**
use **^** and **M-** notation, except for **LFD** and **TAB**
- help** display this help and exit
- version**
output version information and exit

With no **FILE**, or when **FILE** is **-**, read standard input.

AUTHOR

Written by Torbjorn Granlund and Richard M. Stallman.

REPORTING BUGS

Report bugs to [<bug-textutils@gnu.org>](mailto:bug-textutils@gnu.org).

COPYRIGHT

Copyright (C) 2002 Free Software Foundation, Inc.

This is free software; see the source for copying conditions.

There is NO warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

SEE ALSO

The full documentation for cat is maintained as a Texinfo manual. If

the info and cat programs are properly installed at your site, the command

`info cat` should give you access to the complete manual.

Options

- `$ man -f command-name` gives only a brief description of the command
 - `man -f cat`
- `$ whatis command_name` (*equivalent*)

`man -f` and `whatis` display all the occurrences of the command in the chapters
- `$ man -k keyword` displays a list of commands in which description there is the word “keyword”
 - `man -k cat`
 - `man -k manual`

Linux Help Manual

- **info** is a program for reading the documentation about GNU utilities
 - `$ info info` (it shows an introduction to info)
 - `$ info emacs` (it describes how to edit with Emacs)
 - `$ info bash` (it gives a brief description of the Bash shell)
 - `$ info uname` (it prints information about the machine and the operating system it is run on)
 - ✓ `$ uname -a`
Linux homelinux 2.4.19-16mdk #1 Fri Sep 20
18:15:05 CEST 2002 i686 unknown unknown GNU/
Linux

Current Working Directory

- Every process has a location in the directory, termed its current working directory.
- When you log into a Linux system, your shell starts off in a particular directory called home directory (ex. /home/mohamad/).
- To display your current working directory use
 - `$ pwd`
- To come back the home directory use
 - `$ cd`
- Each file or directory has a unambiguously name specified by its pathname relative to / the root directory.
- A pathname relative the root directory is often termed an absolute pathname (ex. /usr/src/linux/CREDITS).
 - `$ cd /usr/src/linux/CREDITS`

Absolute and relative pathname

- A file can be unambiguously specified by using a pathname relative to its current working directory.
 - Example: if `/usr/src/` is the current working directory, `/linux/CREDITS` is the relative pathname of `CREDITS`
- The file system provides the following special fields that may be used supplying a relative path:
 - `.` - current directory
 - `..` - parent directory
 - `~` - your home directory
 - Example:
 - ✓ `cd ..`
 - ✓ `cd ..`
 - ✓ `cd game`
 - ✓ `./fortune`

Listing the contents of a directory

- `$ ls` lists contents of current directory
- `$ ls dir_name` list `dir_name` contents
- **ls options**
 - `-a` all files including hidden files
 - `-C` column list sorted down
 - `-F` adds / for directory; * for executable; @ for symbolic links
 - `-l` long format - file details
 - `-m` across page; comma separators
 - `-p` adds / for directory; * for executable
 - `-r` reverse alphabet order
 - `-R` recursive; includes subdirectories
 - `-s` size of files in blocks
 - `-t` list in time date last modified
 - `-u` lists in time date last accessed
 - `-x` column list sorted across page
 - `-i` inode of each file

Creating and Removing Directories

- `$ mkdir dir_name` (creates a directory)
 - `$ mkdir appunti ; ls`
 - `$ mkdir {appunti, lucidi}; ls`
- `$ rmdir dir_name` (removes empty directory. No warning!)
 - `$ rmdir appunti`
- If the directory is not empty, to remove it we must use `$ rm -r dir_name`
 - `$ rm -r appunti`
- It is not possible to remove the directories between home and /

Creating and Removing Files

- `$ cat > file_name` (store keyboard input into the file)

➤ `$ cat > prova`

Help! I'm stuck in a Linux program!

^D

(The red text indicates what the user types.)

- `$ rm file_name` (removes file. No warning!)

➤ `$ rm pippo` (Permanently removes file pippo. No warning!)

➤ rm options:

- ✓ `-r` (*recursive*): removes the contents of directories recursively
- ✓ `-i` (*interactive*): prompts whether to remove each file
- ✓ `-f` (*force*): forces rm to remove files independently from the permissions

Copy

- `$ cp options file1 file2` (to copy file1 into file2)
 - `$ cp /etc/passwd pass`
 - `$ cp problemi/* ~/backup`
- If `file2` does not exist, then `cp` creates it; otherwise `cp` overwrites it
- If `file2` is a directory, `cp` makes a copy of `file1` in the directory
 - `$ cp pippo ~/articoli`
 - `$ cp /etc/passwd .`
- `cp` options:
 - `-i` prompt before overwriting existing file
 - `-ppreserve` permissions
 - `-r` recursive copy files and subdirectories

Move

- `$ mv olddirectory newdirectory (renames directory oldname to newname)`
 - If `newdirectory` already exists `mv` moves `olddirectory` into the new one
- `$ mv oldname newname (renames file oldname to newname)`
 - If `newname` already exists `mv` writes `oldname` over `newname`
- `mv` options:
 - `-i` prompt before overwriting existing file
 - `-f` forces `mv` to replace reserve permissions
- `$ mv file path (moves file in current directory to new directory)`
 - `$ mv chap[1,3,7] book` (moves files `chap1`, `chap3`, and `chap7` to directory `book`)
 - `$ mv chap[1-5] book` (moves files `chap1` to `chap5` to directory `book`)

Listing a file

- `$ cat filename` (displays the contents of filename)
- `$ more filename` (displays first screen of filename use space bar to scroll up one screen quits automatically after last screen)
- `$ less filename` (displays first screen of filename use space bar to scroll up one screen use up and down arrows to move up or down one line need to type q or Q to exit less command)
- `$ head -n filename` (displays the first n lines of filename. If n is not specified, it defaults to 10)
- `$ tail -n filename` (displays the last n line of filename. If n is not specified, it defaults to 10)

Standard Input and Standard Output

- Every program you run from the shell opens three files:
 - standard input ← 0
 - standard output ← 1
 - standard error ← 2
- The files provide the primary means of communications between the programs, and exist for as long as the process runs.
- **The standard input** file provides a way to send data to a process. As a default, the standard input is read from the *terminal keyboard*.
- **The standard output** provides a means for the program to output data. As a default, the standard output goes to the *terminal display screen*.
- **The standard error** is where the program reports any errors encountered during execution. By default, the standard error goes to the *terminal display*.

Redirecting Input and Output

- It is possible to tell a program:
 - where to look for input
 - where to send output,using input/output redirection. UNIX uses the special characters **<** and **>** to signify input and output redirection, respectively.
- Redirecting input: Using **<** with a file name (i.e., **< file1**) in a shell command tells the shell to read input from a file called "file1" instead of from the keyboard.
 - `$ more < /etc/passwd`
- Redirecting output: Using **>** with a file name (i.e., **> file 2**) causes the shell to place the output from the command in a file called "file2" instead of on the screen. If the file "file2" already exists, the old version will be overwritten.
 - `$ ls /tmp > ~/ls.out`
 - `$ sort pippo > pippo.ordinato`

Redirecting Input and Output

- Use **>>** to append to an existing file (i.e., **>> file2**) causes the shell to append the output from a command to the end of a file called "file2".
 - If the file "file2" does not already exist, it will be created.

- Example

```
1. $ ls /bin > ~/bin; wc -l ~/bin  
$ ls /usr/sbin > ~/bin ; wc -l ~/bin
```

```
2. $ ls /bin > ~/bin;  
$ ls /usr/sbin >> ~/bin; wc -l ~/bin
```


Redirecting Error

- Using **>&** with a file name (i.e., **>& file1**) causes the shell to place the **standard error** and the **standard output** from the command in a file called "file1".
 - If the file "file1" already exists, the old version will be overwritten.
- **Example**
 - `$ ls abcdef`
 - `$ ls abcdef >& lerror`
 - `cat lerror`
 - `$ abcdef >& command`
 - `cat command`
 - `$ mkdir /bin/miei >& ~/miei; cat ~/miei`
 - `$ rm /bin/perl >& ~/errperl; cat ~/errperl`

Pipes

- UNIX allows you to connect processes, by letting the standard output of one process feed into the standard input of another process. That mechanism is called a **pipe (|)**.
- `$ command1 | command2` causes the standard output of `command1` to flow through to standard input of `command2`.
- A sequence of commands chained together in this way is called a **pipeline**. Connecting simple processes in a pipeline allows you to perform complex tasks without writing complex programs.
 - `$ cat /etc/passwd | sort > ~/pass_ord`
 - `$ sort < pippo | lpr`

Exercises

1. Determine the number of files in the directory /bin whose first letter is "c"
2. Create a file containing the names of the first 7 files of the directory /etc
3. Determine the number of files of the directory /etc that contain "."
4. Create a file containing a list with the name of 10 commands of /bin sorting by the last access time
5. Create a file containing the names of the first 7 files and the last 6 files (sorted in alphabetical order) of the directory /etc
6. Create a file containing a list with the name of 8 files in /usr/sbin/ sorting by the last modification time